

CptS 415 / Assignment-05

1. MapReduce

a. Common Friends

Facebook updates the “common friends” of you and respond to hundreds of millions of requests every day.

The friendship information is stored as a pair: `(Person, [List of Friends])` for every user in the social network.

Write a MapReduce program *to return a dictionary of common friends* of the form:

```
user_pair, list_of_common_friends = \
    (user_i, user_j),
    [
        list(
            #Common Friends of user_i and user_j
        )
    ]

# for all pairs of i and j who are friends.
```

Python

The order of i and j you returned should be the same as the lexicographical order of their names.

You need to give the pseudo-code of a main function, and both Map() and Reduce() function. Specify the key/value pair and their semantics (what are they referring to?).

△ Solution

```
class Person:
    def __init__(self, name: str):
        self.name: str = name

class User:
    """
    Each User object contains:
    - k1: a Person, a unique object in database
    - v1: friendship, a dict associated with the primary key, where
        each entry use the address of the person in database as key
        for uniqueness.
    """
    def __init__(self, person: Person, friends: dict[str, Person]):
        self.person: Person = person
        self.friendship: dict[str, Person] = friends
        self.friends_count: int = len(self.friendship)
    def __repr__(self):
        return f"({self.person.name}, {[name for name in self.friendship]})"

def Map(i: User, j: User):
    """
```

Python

```

Compare the two User objects and return a new 2-tuple:
- k2: the shorter friend list as upper bound for potential friends
- v2: the other person as target for matching
"""
if (i.friends_count > j.friends_count):
    return (i.friendship, j)
else:
    return (j.friendship, i)

def Reduce(potential: dict[str, Person], target: User):
    """
    From a list of potential friends and a target, check every
    key in the potential list against the target's friendship.
    Return the list of common keys.
    """
    if len(potential) == 0 or len(target.friendship) == 0:
        return []
    else:
        common: list[Person] = list()
        for person in potential:
            if person in target.friendship:
                common += [potential[person]]
        return common

def MapReduce(i: User, j: User):
    """
    Map and Reduce together.
    - k3: the pair of users
    - v3: a list of Person objects in database
    """
    potential, target = Map(i, j)
    common: list[Person] = Reduce(potential, target)
    user_pair = sorted([i.person.name, j.person.name])
    return {"pair": tuple(user_pair), "common": common}

def main():
    Alice = Person('Alice')
    Bob = Person('Bob')
    John = Person('John')
    Jane = Person('Jane')

    a = User(Person('Jack'), {
        f"{Bob}": Bob,
        f"{Jane}": Jane,
        f"{Alice}": Alice,
        f"{John}": John
    })
    b = User(Person('Mary'), {
        f"{Jane}": Jane,
        f"{Bob}": Bob,
        f"{Alice}": Alice
    })

    result = MapReduce(a, b)
    print(result["pair"], result["common"])

```

```
if __name__ == "__main__":
    main()
```

b. Top-10 Keywords

Search engine companies like Google maintains hot webpages in a set R for keyword search. Each record $r \in R$ is an article, stored as a sequence of keywords. Write a MapReduce program to report the top 10 most frequent keywords appeared in the webpages in R .

Give the pseudo-code of your MR program.

Solution

Python

```
def Map(R: list[list[str]]):
    """
    Map every word in every article r in R to an
    entry in a dictionary:
    - k_w: word
    - v_w: frequency of the word, incremented at current iteration
    - input: list of lists of words
    - output: dictionary of words and frequency
    """
    bag: dict[str, int] = dict()
    for r in R:
        for word in r:
            bag[word] = bag.get(word, 0) + 1
    return bag

def Reduce(bag: dict[str, int]):
    """
    From a bag of words and associated frequencies,
    sort the bag by value (word frequency) and return an ordered map
    of words and their associated frequency.
    input: unordered map
    output: ordered list of 2-tuples
    """
    dict(sorted(bag.items(), key=lambda item: item[1]))
    return list(bag.items())

def MapReduce(R: list[list[str]]):
    """
    return keys from the last 10 entries as most frequent words.
    """
    word_frequency_map = Map(R)
    words_ordered_list = Reduce(word_frequency_map)
    list_size = len(words_ordered_list)
    return words_ordered_list[list_size-10:-1]
```

2. Graph Parallel Models: MR for Graph Processing

a.

Consider the common friends problem in Problem 1.a. We study a “2-hop common contact problem”, where a list should be returned for any pair of friends i and j , such that the list contains all the users that can reach both i and j within 2 hops. Write a MR algorithm to solve the problem and give the pseudo code.

Solution

b.

We described how to compute distances with mapReduce. Consider a class of d -bounded reachability queries as follows. Given a graph G , two nodes u and v and an integer d , it returns a Boolean answer **YES**, if the two nodes can be connected by a path of length no greater than d . Otherwise, it returns **NO**. Write an MR program to compute the query $Q(G, u, v, d)$ and give the pseudo code.

Provide necessary correctness and complexity analysis.

Solution

3. Hadoop

Hadoop Program:

The attached CSV file contains hourly normal recordings for temperature and dew point temperature at Asheville Regional Airport, NC, USA. *The unit of measurement* is in **tenths of a degree Fahrenheit**. For example, 344 is 34.4 F.

Write a program using Hadoop to compute and output daily average measurements for temperature and dew point temperature.

The daily average measurements *should include measurements for 24-hour period*. For example, from:

```
20100101 00:00 (2010, January 1st, 00:00)
```

to:

```
20100101 23:00 (2010, January 1st, 23:00)
```

Output the result in the format shown below - the columns are date and the combined result (separated by comma) of daily temperature and daily dew point temperature:

```
20100101    377.04, 285.58
20100102    378.67, 286.92
....       .... , ....
```

Plain Text

You may write the application in Java, C/C++ or Python language. Provide both source code and compiled code, if applicable, for your program.

Solution

- First, I need to look up the API for Hadoop in Python/C++.